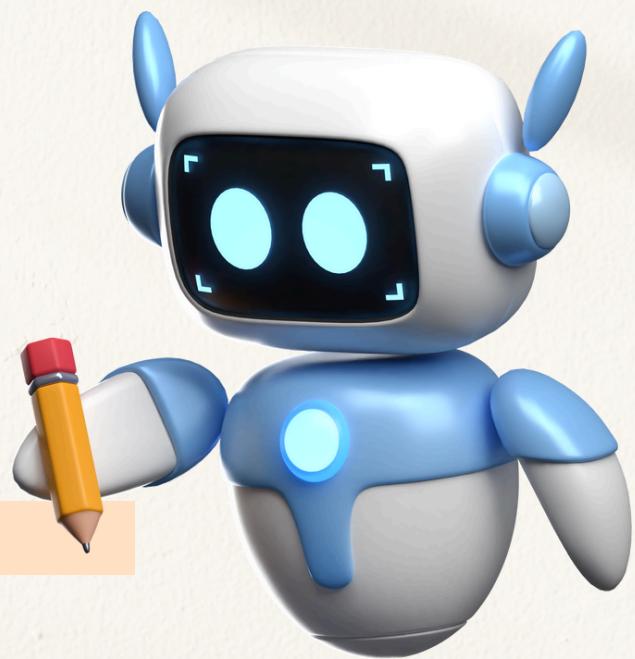


#4

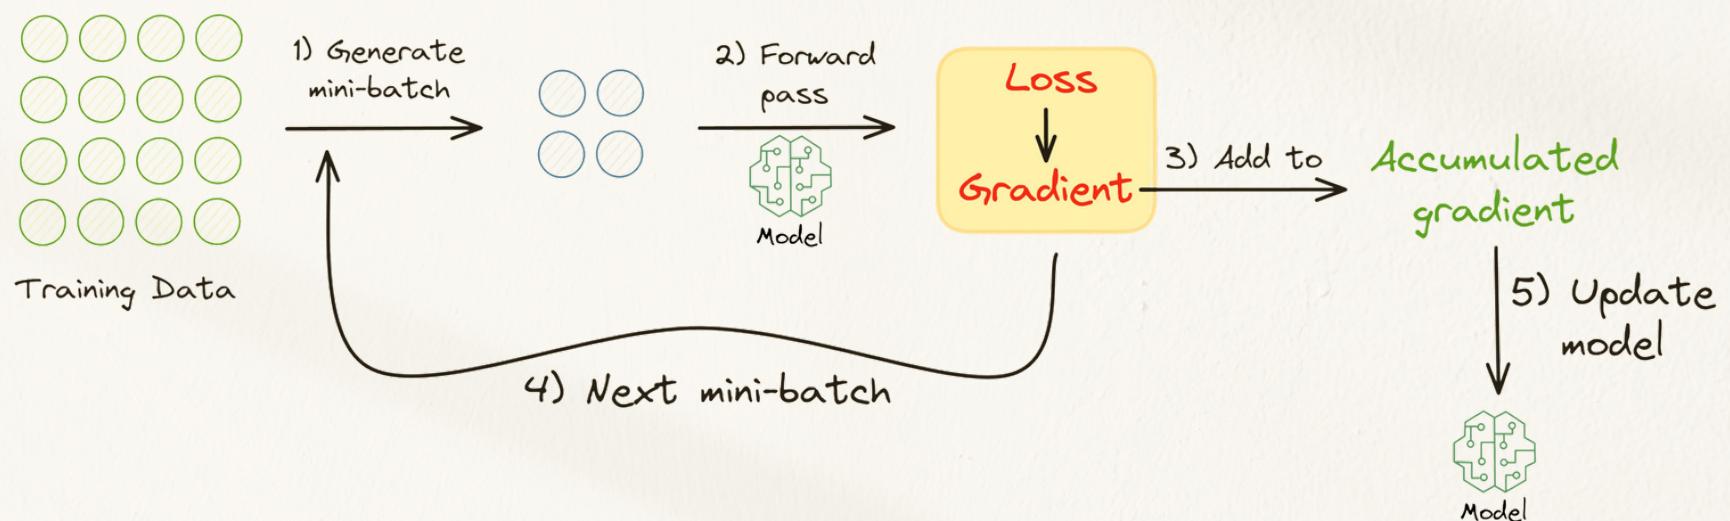
INTERVIEW QUESTION

What is gradient accumulation, and how does it help when training large models?

→
SWIPE NEXT

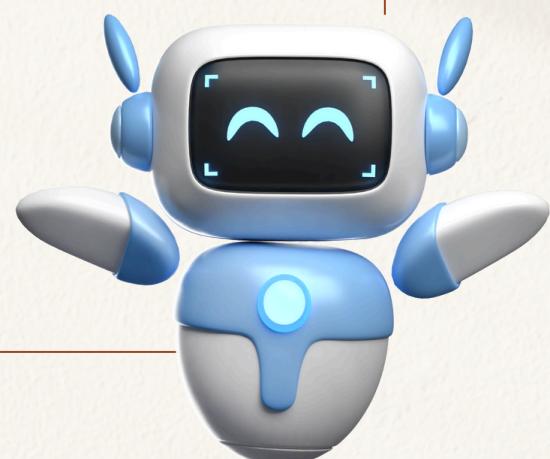
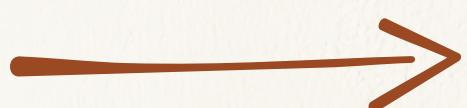


When you have limited memory, gradient accumulation allows you to simulate a larger batch size by accumulating gradients over several smaller batches before updating the model. This ensures stable training even when memory constraints prevent large batch sizes.



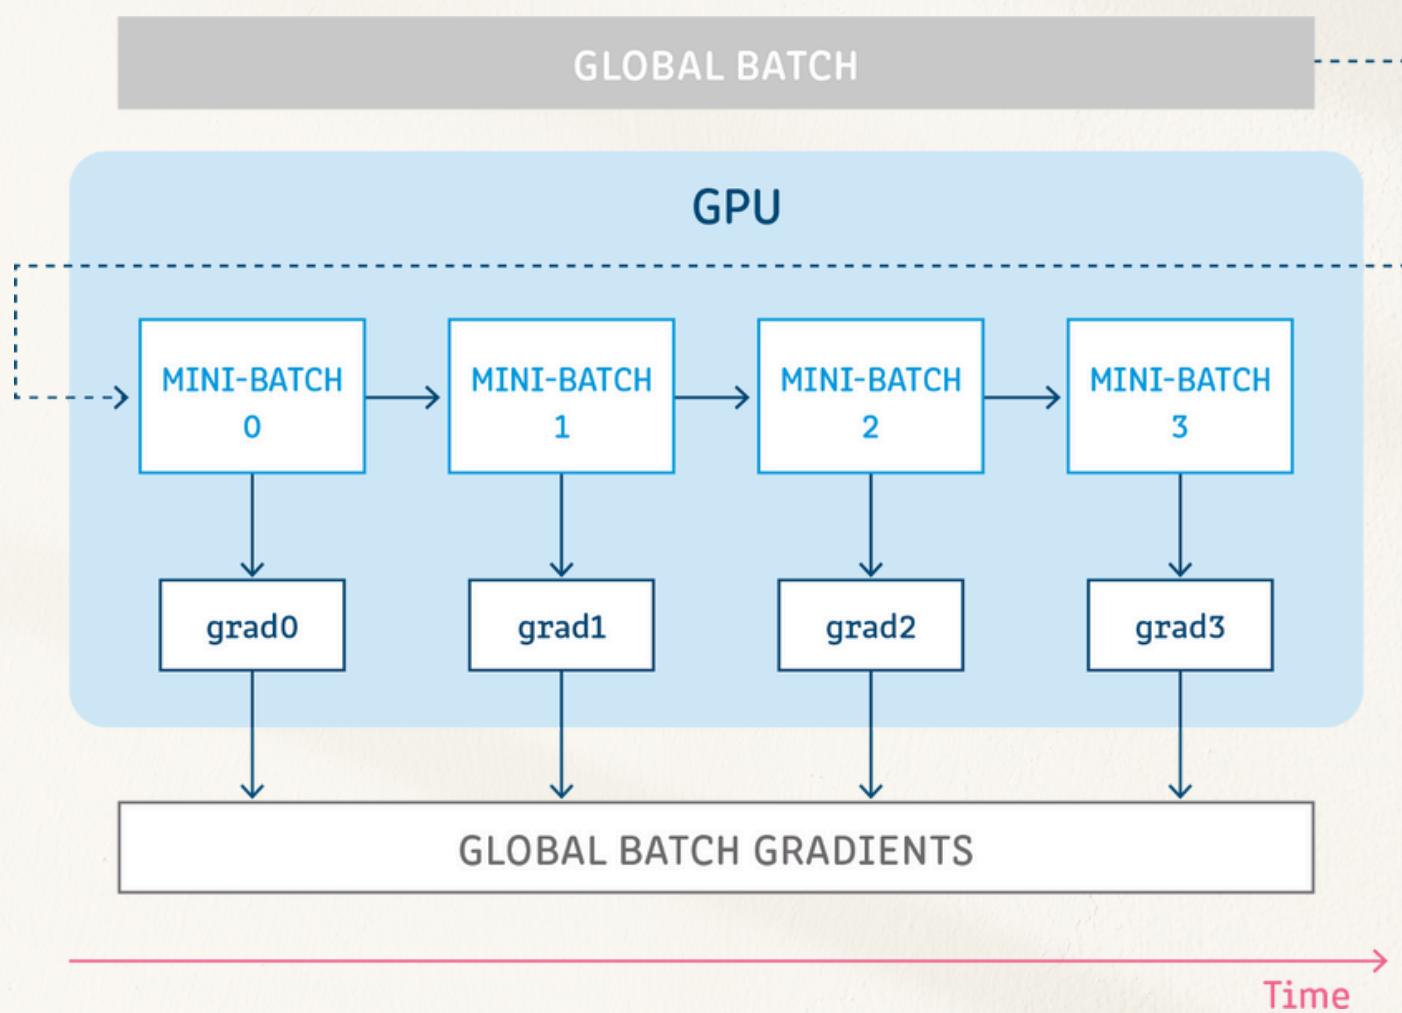
SEE THE IMAGE

ON NEXT PAGE

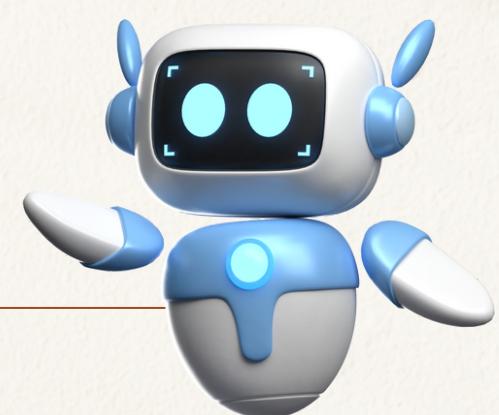


KEY LEARNING

Gradient accumulation allows for flexible batch sizes without compromising on training quality.



CODE Snippet

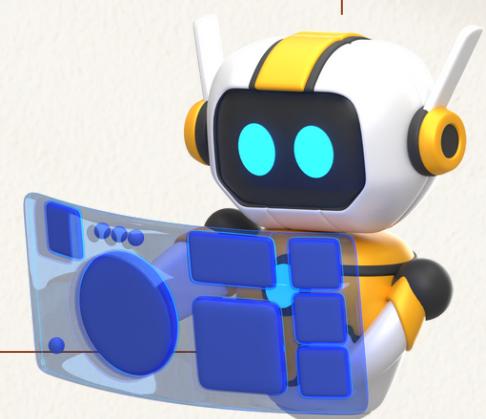


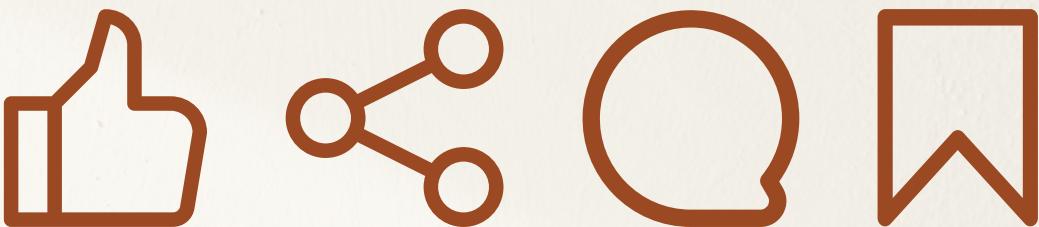
<CODE/>

```
for inputs, labels in trainloader:  
    # Clear old gradients  
    optimizer.zero_grad()  
    # Run forward pass  
    outputs = net(inputs)  
    # Compute Loss  
    loss = criterion(outputs, labels)  
    # Compute gradients  
    loss.backward()  
    # Update weights  
    optimizer.step()
```

WITHOUT GRADIENT ACCUMULATION

```
accumulation_step = 5  
for idx, (inputs, labels) in enumerate(trainloader):  
    # Run forward pass  
    outputs = net(inputs)  
    # Compute Loss  
    loss = criterion(outputs, labels)  
    # Compute gradients  
    loss.backward()  
    # Update weights every few iterations and check if it is last mini batch  
    if ( (idx + 1) % accumulation_step == 0) or ((idx + 1) == len(trainloader)):  
        # Update weights  
        optimizer.step()  
        # Clear old gradients  
        optimizer.zero_grad()
```

WITH GRADIENT ACCUMULATION



**DON'T
FORGET TO
LIKE, SHARE
COMMENT
AND SAVE**

